

****HW1 - 15 Points****

- **You have to submit two files for this part of the HW**

(1) ipynb (colab notebook) and
(2) pdf file (pdf version of the colab file).**

- **Files should be named as follows:**

FirstName_LastName_HW_1**

```
In [1]: import torch
import time
```

Q1 : Create Tensor (1/2 Point)

Create a torch Tensor of shape (5, 3) which is filled with zeros. Modify the tensor to set element (0, 2) to 10 and element (2, 0) to 100.

```
In [2]: my_tensor = torch.zeros(5, 3)
```

```
In [3]: my_tensor.shape
```

```
Out[3]: torch.Size([5, 3])
```

```
In [4]: my_tensor
```

```
Out[4]: tensor([[0., 0., 0.],
               [0., 0., 0.],
               [0., 0., 0.],
               [0., 0., 0.],
               [0., 0., 0.]])
```

```
In [5]: # Manually set the value at the first row and third column to 10,
# and the value at the third row and first column to 100 in the tensor named "my_tensor"

my_tensor[0, 2] = 10
my_tensor[2, 0] = 100
```

```
In [6]: my_tensor
```

```
Out[6]: tensor([[ 0.,  0., 10.],
               [ 0.,  0.,  0.],
               [100.,  0.,  0.],
               [ 0.,  0.,  0.],
               [ 0.,  0.,  0.]])
```

Q2: Reshape tensor (1/2 Point)

You have following tensor as input:

```
x=torch.tensor([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23])
```

Using only reshaping functions (like view, reshape, transpose, permute), you need to get at the following tensor as output:

```
tensor([[ 0,  4,  8, 12, 16, 20],
        [ 1,  5,  9, 13, 17, 21],
        [ 2,  6, 10, 14, 18, 22],
        [ 3,  7, 11, 15, 19, 23]])
```

```
In [7]: x=torch.tensor([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23])
```

```
In [8]: x = x.view(6, 4).transpose(0,1)
x
```

```
Out[8]: tensor([[ 0,  4,  8, 12, 16, 20],
               [ 1,  5,  9, 13, 17, 21],
               [ 2,  6, 10, 14, 18, 22],
               [ 3,  7, 11, 15, 19, 23]])
```

Q3: Slice tensor (1Point)

- Slice the tensor x to get the following

- last row of x
- fourth column of x
- first three rows and first two columns - the shape of subtensor should be (3,2)
- odd valued rows and columns

```
In [9]: x = torch.tensor([[1, 2, 3, 4, 5], [6, 7, 8, 8, 10], [11, 12, 13, 14, 15]])
x
```

```
Out[9]: tensor([[ 1,  2,  3,  4,  5],
               [ 6,  7,  8,  8, 10],
               [11, 12, 13, 14, 15]])
```

```
In [10]: x.shape
```

```
Out[10]: torch.Size([3, 5])
```

```
In [11]: # Student Task: Retrieve the last row of the tensor 'x'
# Hint: Negative indexing can help you select rows or columns counting from the end.
# Think about how you can select all columns for the desired row.
last_row = x[-1, :]
last_row
```

```
Out[11]: tensor([11, 12, 13, 14, 15])
```

```
In [12]: # Student Task: Retrieve the fourth column of the tensor 'x'
# Hint: Pay attention to the indexing for both rows and columns.
# Remember that indexing in Python starts from zero.
fourth_column = x[:, 3]
fourth_column
```

```
Out[12]: tensor([ 4,  8, 14])
```

```
In [13]: # Student Task: Retrieve the first 3 rows and first 2 columns from the tensor
# Hint: Use slicing to extract the required subset of rows and columns.
first_3_rows_2_columns = x[:3, :2]
first_3_rows_2_columns
```

```
Out[13]: tensor([[ 1,  2],
                 [ 6,  7],
                 [11, 12]])
```

```
In [14]: # Student Task: Retrieve the rows and columns with odd-indexed positions from
# Hint: Use stride slicing to extract the required subset of rows and columns
odd_valued_rows_columns = x[::2, ::2]
odd_valued_rows_columns
```

```
Out[14]: tensor([[ 1,  3,  5],
                 [11, 13, 15]])
```

Q4 -Normalize Function (1/2 Points)

Write the function that normalizes the columns of a matrix. You have to compute the mean and standard deviation of each column. Then for each element of the column, you subtract the mean and divide by the standard deviation.

```
In [15]: # Given Data
x = [[ 3,  60,  100, -100],
      [ 2,  20,  600, -600],
      [-5,  50,  900, -900]]
```

```
In [16]: # Convert to PyTorch Tensor and set to float
X = torch.tensor(x)
X = X.float()
```

```
In [17]: # Print shape and data type for verification
print(X.shape)
print(X.dtype)

torch.Size([3, 4])
torch.float32
```

```
In [18]: # Compute and display the mean and standard deviation of each column for reference
X.mean(axis = 0)
X.std(axis = 0)
```

```
Out[18]: tensor([ 4.3589, 20.8167, 404.1452, 404.1452])
```

```
In [19]: X.std(axis = 0)
```

```
Out[19]: tensor([ 4.3589, 20.8167, 404.1452, 404.1452])
```

- Your task starts here
- Your `normalize_matrix` function should take a PyTorch tensor `x` as input.
- It should return a tensor where the columns are normalized.
- After implementing your function, use the code provided to verify if the mean for each column in `Z` is close to zero and the standard deviation is 1.

```
In [20]: def normalize_matrix(x):  
# Calculate the mean along each column (think carefully , you will take mean  
mean = torch.mean(x, dim=0)  
  
# Calculate the standard deviation along each column  
std = torch.std(x, dim=0)  
  
# Normalize each element in the columns by subtracting the mean and dividing  
y = (x - mean) / std  
  
return y # Return the normalized matrix
```

```
In [21]: Z = normalize_matrix(X)  
Z
```

```
Out[21]: tensor([[ 0.6882,  0.8006, -1.0722,  1.0722],  
[ 0.4588, -1.1209,  0.1650, -0.1650],  
[-1.1471,  0.3203,  0.9073, -0.9073]])
```

```
In [22]: Z.mean(axis = 0)
```

```
Out[22]: tensor([ 0.0000e+00,  4.9671e-08,  3.9736e-08, -3.9736e-08])
```

Q5: In-place vs. Out-of-place Operations (1 Point)

1. Create a tensor `A` with values `[1, 2, 3]`.
2. Perform an in-place addition (use `add_` method) of `5` to tensor `A`.
3. Then, create another tensor `B` with values `[4, 5, 6]` and perform an out-of-place addition of `5`.

Print the memory addresses of `A` and `B` before and after the operations to demonstrate the difference in memory usage. Provide explanation

```
In [23]: A = torch.tensor([1, 2, 3])  
print('Original memory address of A:', id(A))  
A.add_(5)  
print('Memory address of A after in-place addition:', id(A))  
print('A after in-place addition:', A)  
  
B = torch.tensor([4, 5, 6])  
print('Original memory address of B:', id(B))
```

```
B = B + 5
print('Memory address of B after out-of-place addition:', id(B))
print('B after out-of-place addition:', B)
```

Original memory address of A: 5434774480
 Memory address of A after in-place addition: 5434774480
 A after in-place addition: tensor([6, 7, 8])
 Original memory address of B: 5434774288
 Memory address of B after out-of-place addition: 5434773040
 B after out-of-place addition: tensor([9, 10, 11])

Provide Explanation for above question here :

A is a in place operation. This ensures memory address of the tensor remains the same and helps with memory allocation.

B is being created again and assigned the same variable name. Therefore, each time B is created it is allotted a new address.

Q6: Tensor Broadcasting (1 Point)

1. Create two tensors `X` with shape `(3, 1)` and `Y` with shape `(1, 3)`. Perform an addition operation on `X` and `Y`.
2. Explain how broadcasting is applied in this operation by describing the shape changes that occur internally.

```
In [24]: X = torch.rand((3, 1))
Y = torch.rand((1, 3))
print('Original shapes:', X.shape, Y.shape)
print(X)
print(Y)
result = X + Y
print('Result:', result)
print('Result shape:', result.shape)
```

Original shapes: torch.Size([3, 1]) torch.Size([1, 3])
 tensor([[0.0992],
 [0.1627],
 [0.6817]])
 tensor([[0.3329, 0.4364, 0.1186]])
 Result: tensor([[0.4321, 0.5356, 0.2178],
 [0.4956, 0.5991, 0.2813],
 [1.0146, 1.1180, 0.8002]])
 Result shape: torch.Size([3, 3])

Provide Explanation for above question here :

Broadcasting enables the tensors X and Y to be of same shape and it used to perform the Tensor addition

- PyTorch identifies compatible dimensions (here, the first dimension of X and the second dimension of Y have a size of 1).

- It virtually expands those dimensions to match each other, creating intermediate views of X and Y as if they both had a shape of (3, 3).
- This allows element-wise addition to proceed without explicit reshaping.

Q7: Linear Algebra Operations (1 Point)

1. Create two matrices `M1` and `M2` of compatible shapes for matrix multiplication. Perform the multiplication and print the result.
2. Then, create two vectors `V1` and `V2` and compute their dot product.

```
In [25]: M1 = torch.ones(3,2)
M2 = torch.ones(2,3)
mat_multiplication = torch.mm(M1,M2)
print('Matrix multiplication result:', mat_multiplication)

V1 = torch.ones(3,)
V2 = torch.ones(3,)
dot_product = torch.dot(V1, V2)
print('Dot product:', dot_product)
```

Matrix multiplication result: tensor([[2., 2., 2.],
[2., 2., 2.],
[2., 2., 2.]])
Dot product: tensor(3.)

Q8: Manipulating Tensor Shapes (1 Point)

Given a tensor `T` with shape `(2, 3, 4)`, demonstrate how to

1. reshape it to `(3, 8)` using view,
2. reshape it to `(4, 2, 3)` using reshape,
3. transpose the first and last dimensions using permute.
4. explain what is the difference between reshape and view

```
In [26]: T = torch.rand(2, 3, 4)
T_view = T.view(3,8)
print('T_view shape:', T_view.shape)

T_reshape = T.reshape(4, 2, 3)
print('T_reshape shape:', T_reshape.shape)

T_permute = T.permute(2, 1, 0)
print('T_permute shape:', T_permute.shape)
```

T_view shape: torch.Size([3, 8])
T_reshape shape: torch.Size([4, 2, 3])
T_permute shape: torch.Size([4, 3, 2])

Provide Explanation for above question here :

View can be only used for contiguous tensor in contrary to reshape.

If you just want to reshape tensors, use `torch.reshape`. If you're also concerned about memory usage and want to ensure that the two tensors share the same data, use `torch.view`.

Q9: Tensor Concatenation and Stacking (1 Point)

Create tensors `C1` and `C2` both with shape (2, 3).

1. Concatenate them along dimension 0 and then along dimension 1. Print the shape of the resulting tensor.
2. Afterwards, stack the same tensors along dimension 0 and print the shape of the resulting tensor.
3. What is the difference between stacking and concatenating.

```
In [27]: C1 = torch.rand(2, 3)
C2 = torch.rand(2, 3)
concatenated_dim0 = torch.cat([C1, C2], dim=0)
print('Concatenated along dimension 0:', concatenated_dim0.shape)

concatenated_dim1 = torch.cat([C1, C2], dim=1)
print('Concatenated along dimension 1:', concatenated_dim1.shape)

stacked = torch.stack([C1, C2], dim=0)
print('Stacked tensor shape:', stacked.shape)
```

```
Concatenated along dimension 0: torch.Size([4, 3])
Concatenated along dimension 1: torch.Size([2, 6])
Stacked tensor shape: torch.Size([2, 2, 3])
```

Explain the difference between concatenating and stacking here

Concatenating combines tensors along an existing dimension, whereas stacking adds a dimension and adds the tensors along that dimension.

Q10: Advanced Indexing and Slicing (1 Point)

1. Given a tensor `D` with shape (6, 6), extract elements that are greater than 0.5.
2. Then, extract the second and fourth rows from `D`.
3. Finally, extract a sub-tensor from the top-left 3x3 block.

```
In [28]: D = torch.rand(6, 6)
print('Elements greater than 0.5:\n', D[D > 0.5])

second_fourth_rows = D[[1, 3], :]
```

```
print('\nSecond and fourth rows:\n', second_fourth_rows)
```

```
top_left_3x3 = D[:3, :3]  
print('\nTop-left 3x3 block:\n ', top_left_3x3)
```

Elements greater than 0.5:

```
tensor([0.8305, 0.5064, 0.9334, 0.8831, 0.8864, 0.8975, 0.5566, 0.9483, 0.8766,  
        0.5091, 0.9582, 0.6166, 0.8757, 0.9742, 0.7940, 0.6839, 0.8771, 0.7771])
```

Second and fourth rows:

```
tensor([[0.1285, 0.0165, 0.3430, 0.0222, 0.3897, 0.3624],  
        [0.8766, 0.5091, 0.9582, 0.6166, 0.8757, 0.1016]])
```

Top-left 3x3 block:

```
tensor([[0.2052, 0.8305, 0.5064],  
        [0.1285, 0.0165, 0.3430],  
        [0.1820, 0.8831, 0.8864]])
```

Q11: Tensor Mathematical Operations (1 Point)

1. Create a tensor `G` with values from 0 to π in steps of $\pi/4$.
2. Compute and print the sine, cosine, and tangent logarithm and the exponential of `G`.

```
In [29]: import numpy as np  
G = torch.arange(0, torch.tensor(np.pi), step=torch.tensor(np.pi / 4))  
  
print('G:', G)  
print('Sine of G:', torch.sin(G))  
print('Cosine of G:', torch.cos(G))  
print('Tangent of G:', torch.tan(G))  
print('Natural logarithm of G:', torch.log(G))  
print('Exponential of G:', torch.exp(G))
```

```
G: tensor([0.0000, 0.7854, 1.5708, 2.3562])  
Sine of G: tensor([0.0000, 0.7071, 1.0000, 0.7071])  
Cosine of G: tensor([ 1.0000e+00,  7.0711e-01, -4.3711e-08, -7.0711e-01])  
Tangent of G: tensor([ 0.0000e+00,  1.0000e+00, -2.2877e+07, -1.0000e+00])  
Natural logarithm of G: tensor([ -inf, -0.2416,  0.4516,  0.8570])  
Exponential of G: tensor([ 1.0000,  2.1933,  4.8105, 10.5507])
```

Q12: Tensor Reduction Operations (1 Point)

1. Create a 3x2 tensor `H`.
2. Compute the sum of `H`. Print the result and shape after taking sun.
3. Then, perform the same operations along dimension 0 and dimension 1, printing the results and shapes.
4. What do you observe? How the shape changes?


```
In [30]: H = torch.rand(3, 2)
print('H:', H, end = "\n\n")
print('Shape of original Tensor H', H.shape, end = "\n\n")

print('Sum of H:', torch.sum(H))
print('Shape after Sum of H:', torch.sum(H).shape, end = "\n\n")

print('Sum of H along dimension 0:', torch.sum(H, dim = 0))
print('Shape after sum of H along dimension 0:', torch.sum(H, dim = 0).shape,

print('Sum of H along dimension 1:', torch.sum(H, dim = 1))
print('Shape after sum of H along dimension 1:', torch.sum(H, dim = 1).shape)
```

```
H: tensor([[0.2863, 0.8445],
          [0.3185, 0.5585],
          [0.8606, 0.2818]])
```

```
Shape of original Tensor H torch.Size([3, 2])
```

```
Sum of H: tensor(3.1502)
Shape after Sum of H: torch.Size([])
```

```
Sum of H along dimension 0: tensor([1.4654, 1.6848])
Shape after sum of H along dimension 0: torch.Size([2])
```

```
Sum of H along dimension 1: tensor([1.1308, 0.8770, 1.1424])
Shape after sum of H along dimension 1: torch.Size([3])
```

Provide your observations on shape changes here

1. Is the sum of all the scalars in the Tensor H, therefore it is a scalar.
2. Is the sum along dimension 0 (rows), therefore it is a scalar of dimension (2,).
3. Is the sum along dimension 1 (columns), therefore it is a scalar of dimension (3,).

Q13: Working with Tensor Data Types (1 Point)

1. Create a tensor **I** of data type float with values `[1.0, 2.0, 3.0]`.
2. Convert **I** to data type int and print the result.
3. Explain in which scenarios it's necessary to be cautious about the data type of tensors.

```
In [31]: # Solution for Q16
I = torch.tensor([1.0, 2.0, 3.0], dtype=torch.float)
print('I:', I)
I_int = I.to(dtype=torch.int)
print('I converted to int:', I_int)
```

```
I: tensor([1., 2., 3.])
I converted to int: tensor([1, 2, 3], dtype=torch.int32)
```

Your explanations here

Memory usage, precision loss and gradient computation are the reasons to be cautious.

****Q14. Speedtest for vectorization 1.5 Points****

Your goal is to measure the speed of linear algebra operations for different levels of vectorization.

1. Construct two matrices A and B with Gaussian random entries of size 1024×1024 .
2. Compute $C = AB$ using matrix-matrix operations and report the time. (Hint: Use `torch.mm`)
3. Compute $C = AB$, treating A as a matrix but computing the result for each column of B one at a time. Report the time. (hint use `torch.mv` inside a for loop)
4. Compute $C = AB$, treating A and B as collections of vectors. Report the time. (Hint: use `torch.dot` inside nested for loop)

```
In [32]: ## Solution 1
torch.manual_seed(42) # do not change this
A = torch.randn(1024, 1024)
B = torch.randn(1024, 1024)
```

```
In [33]: ## Solution 2
start = time.time()

C = torch.mm(A, B)

print("Matrix by matrix: " + str(time.time()-start) + " seconds")
```

Matrix by matrix: 0.0035250186920166016 seconds

```
In [34]: ## Solution 3
C = torch.empty(1024,1024)
start = time.time()

C = torch.zeros(1024, 1024)
for i in range(B.size(1)):
    C[:, i] = torch.mv(A, B[:, i])

print("Matrix by vector: " + str(time.time()-start) + " seconds")
```

Matrix by vector: 0.06011795997619629 seconds

```
In [35]: ## Solution 4
C = torch.empty(1024,1024)
start = time.time()

C = torch.zeros(1024, 1024)
for i in range(A.size(1)):
    for j in range(B.size(1)):
        C[i, j] = torch.dot(A[:, i], B[:, j])

print("vector by vector: " + str(time.time()-start) + " seconds")
```

vector by vector: 11.896167039871216 seconds

Q15 : Redo Question 14 by using GPU - 1.5 Points

Using GPUs

How to use GPUs in Google Colab

In Google Colab -- Go to Runtime Tab at top -- select change runtime type -- for hardware accelartor choose GPU

```
In [36]: # Check if GPU is available
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(device)
```

cpu

```
In [37]: ## Solution 1
torch.manual_seed(42)
A= torch.randn((1024, 1024),device=device)
B= torch.randn((1024, 1024),device=device)
```

```
In [38]: print(A)
print(B)
```

```
tensor([[ 1.9269,  1.4873,  0.9007, ...,  1.1085,  0.5544,  1.5818],
        [-1.2248,  0.9629, -1.5785, ..., -0.0334, -0.8276, -0.3524],
        [-0.6002, -0.0580,  0.2975, ..., -2.3891,  0.7178, -1.5831],
        ...,
        [-0.1277,  0.2762,  0.9963, ...,  0.5773,  0.4467,  1.6042],
        [-1.2248, -0.0074, -1.2422, ..., -0.0288,  1.1666, -0.2757],
        [-1.3839, -0.2788, -0.7896, ..., -1.0240, -0.4006,  0.2859]])
tensor([[ -0.6495,  0.6962,  2.4458, ...,  0.8676, -0.7841,  0.4688],
        [-1.8024,  1.1675, -1.0891, ...,  1.8599, -0.6542, -1.1773],
        [-0.0060, -1.6660,  0.4924, ..., -0.1280, -0.6696,  0.8153],
        ...,
        [ 0.8675,  2.4424, -1.4791, ..., -1.4814,  1.6032, -0.7838],
        [ 0.8459,  1.0714, -0.3842, ..., -1.2028, -0.5702,  0.1948],
        [-0.3129,  0.3797, -1.1664, ..., -0.0991, -2.6005,  0.1010]])
```

```
In [39]: ## Solution 2
start=time.time()

C = torch.mm(A, B)

print("Matrix by matrix: " + str(time.time()-start) + " seconds")
```

Matrix by matrix: 0.003468036651611328 seconds

```
In [40]: ## Solution 3
C= torch.empty(1024,1024, device = device)
start = time.time()

C = torch.zeros(1024, 1024)
```

```
for i in range(B.size(1)):
    C[:, i] = torch.mv(A, B[:, i])

print("Matrix by vector: " + str(time.time()-start) + " seconds")
```

Matrix by vector: 0.05757904052734375 seconds

```
In [41]: ## Solution 4
C = torch.empty(1024, 1024, device = device)
start = time.time()

C = torch.zeros(1024, 1024)
for i in range(A.size(1)):
    for j in range(B.size(1)):
        C[i, j] = torch.dot(A[:, i], B[:, j])

print("vector by vector: " + str(time.time()-start) + " seconds")
```

vector by vector: 12.009110689163208 seconds